

Sai Vasanth Vinay Kudari - 50365626

Part - A

Q1: Explain overfitting and underfitting. What evaluation metrics do we use to check if the regression model fits the data well?

Overfitting: A model which overfits has low generalization performance, which means it performs well on training data and its performance is poor on unseen data. In other words, we can say it has high variance and low bias;

Underfitting: A model which underfits performs poorly on both training data and unseen data, which can mean that the model is too simple to fit the data. We can say the model has high bias and low variance.

Some of the commonly used evaluation metrics:

- MAE - Mean Absolute Error
- MSE - Mean Squared Error
- RMSE - Root Mean Squared Error
- RMSLE - Root Mean Squared Log Error
- R-Squared & Adjusted R-Squared

Q2: Differentiate Correlation and Regression. What is R-squared? Discuss the drawback of R-squared in the case of multiple linear regression. What is Adjusted R-squared?

Correlation measures the relationship between two variables, whereas regression tells us how one variable will affect the other. Generally, in regression, we represent the effect of a variable using a line equation, unlike correlation which outputs a single metric.

R-squared: It is an evaluation metric that gives the proportion of variation in the target variable explained by the linear regression model. It is obtained by dividing TSS - RSS with TSS. Where TSS is Total Sum of Squares and RSS is Residual Sum of Squares. If the regression line is very close to the actual points, this means the independent variables explain the majority of variation in the target variable. In such a case, we would have a really high R-squared value.

Drawbacks in case of regression with multiple independent variables

- R-squared value never decreases with addition of redundant independent variables, which might be a problem as they might not add any value to our model but this metric cannot point out this issue.

Adjusted R-squared solves the above problem with multiple linear regression by considering independent variables used for predicting the target variable.

Q3: Explain precision and recall using the confusion matrix. Discuss with examples when precision and recall may consider as more useful classification evaluation metrics over accuracy.

When we are modelling to predict a rare occurrence, one naive model which can give great results by simply predicting -ve as this occurrence is very frequent our model's accuracy will be pretty high but we know this isn't helpful. Examples where precision and recall are more useful

- Predicting a rare disease such as cancer
- Predicting if a person can get an insurance

Evaluation Metrics :

		Actual	
		Correct	Not Correct
Predicted	Selected	True +ve	False +ve
	Not Selected	False -ve	True -ve

Accuracy = $\frac{(True +ve) + (True -ve)}{(True +ve) + (False +ve) + (False -ve) + (True -ve)}$ } When equally distributed

Precision = % of predicted items that are correctly predicted
 $= \frac{True +ve}{(True +ve) + (False +ve)}$

Recall = % of correct items that are predicted
 $= \frac{(True +ve)}{(True +ve) + (False -ve)}$

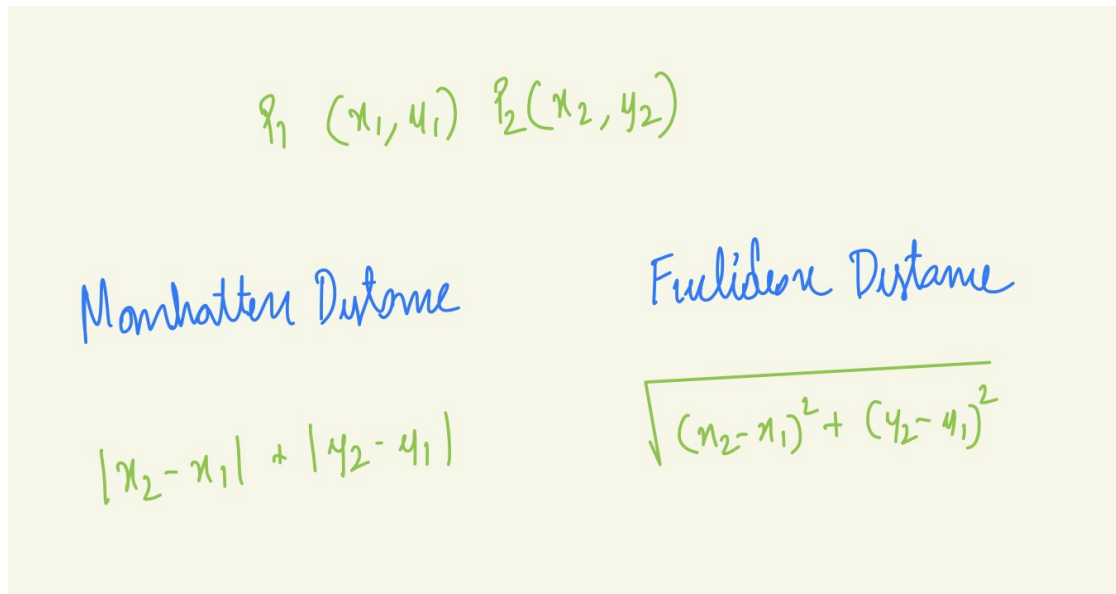
Q4: What is the advantage of K-Means Clustering over Hierarchical Clustering? What are Manhattan Distance and Euclidean Distance in Clustering? Why may we consider different distance measures for clustering?

Both K-Means and Hierarchical clustering are unsupervised algorithms, which aim to recognize the patterns among data points and cluster them, there are couple of advantages of K-Means over Hierarchical clustering

- K-Means is computationally faster than Hierarchical clustering, especially when no of classes are less. Hence, K-Means is suitable for large datasets
- K-Means produces tighter clusters than Hierarchical clustering
- K-means is linear in the number of data objects i.e. $O(n)$, where n is the number of data objects. The time complexity of most of the hierarchical clustering algorithms is quadratic i.e. $O(n^2)$

The choice of distance measure is very important in clustering, this metric will tell us how similar two elements and will influence the shape of clusters. Depending on the type of data

we need to select the distance measure, for instance if two points are similar like most variables but differ on one, Euclidean distance will exaggerate the difference, but Manhattan distance would ignore it, being more influenced by the closeness of the other variables.



The image shows handwritten mathematical formulas for two distance metrics. At the top, two points are labeled: $p_1(x_1, y_1)$ and $p_2(x_2, y_2)$. Below this, the 'Manhattan Distance' is defined as the sum of absolute differences: $|x_2 - x_1| + |y_2 - y_1|$. To the right, the 'Euclidean Distance' is defined as the square root of the sum of squared differences: $\sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$.

Manhattan distance is generally preferred over Euclidean when there is high dimensionality in the data

Q5: What is feature scaling? Show an example of why feature scaling is crucial? What are the two standard feature scaling techniques?

Learning algorithms depend on feature values to predict targets, often the features values are in different scales for instance in modelling country GDP the features may include per capita income, population, no of people exceeding certain income level, in this case we can understand that scales are quite different and there might be a chance that our model might give more importance to certain features. Hence, we need to scale the features before feeding into learning algorithm.

Example: In case of K-Means clustering we use some distance metric such as Euclidean/Manhattan distance to find similarity between two data elements, if we dont scale our distance metric could give wrong results when its comparing similarity between two differently scaled elements such as population vector and per capita income vector.

Feature Scaling Techniques:

- Normalization is a scaling technique that shifts and rescales numbers so that they end up ranging between 0 and 1. It is also known as Min-Max scaling.
- Standardization is another scaling strategy, in which the values are centered around the mean and have a unit standard deviation. This means that the attribute's mean becomes zero, and the resulting distribution has a standard deviation of one.

Q6: Discuss collaborative and content-based recommendation techniques with examples

Collaborative recommendations are based on closely related entities preferences, for instance in movie recommendation engine, it will find users similar to the current user and it recommends the movies which are watched/liked by similar users. This type of recommendation works well but it also has a drawback, it has the cold start problem, it cannot work well unless we have some users already in system, which is not in case of content based recommendation systems.

Content based ones solely depend on items and its features to recommend to users like taking preference from other closely associated users. This engine will leverage user item interactions to learn the user's preference and model recommendations accordingly, it suffers far less from cold start problem but it will have a problem recommending new items which are contrasting to users preferences but they may be liked by the user.

Often in practice, state of the art recommendation engines use both content and collaborative engines in tandem

Part - B

Traffic Violation Analysis

After correlation analysis on dataset I found 'Description', 'Contributed To Accident', 'Belts', 'Personal Injury', 'Property Damage' are useful in predicting target accurately. Alongside I've engineered a new feature 'Description Pred' which is result of text classification on 'Description' column which seems to be highly correlated as expected.

All the models do a decent job at classifying and perform equally good, but Random Forest Classifier is slightly better when compared to other two, it consistently gave better accuracy around 80% and also good sensitivity and specificity scores. Sensitivity score for 'Citation' class using Random Forest and KNN is around 0.90 and using SVM it is around 0.80. We can also observe that 'SERO' class has higher sensitivity (close to 1) and specificity among all other models which is due class imbalance, number of samples with SERO class are the least in the dataset and the descriptions are quite similar. Also, our the models have tough time predicting 'Warning' class this can be concluded by lower sensitivity scores.

Among all the three models Random Forest gives better accuracy but I found KNN more generalizable even with slightly lower accuracy.

I would deploy KNN in production for two reasons

- Inference time is faster than other models : In production we should care about faster inferences times as we need to cater for efficiency and scale
- Lower generalization gap : KNN and SVM both generalize well, which mean that training the test error gap is lower, but since KNN has both accuracy and low generalization gap, I'm inclined toward KNN

Please find below charts and scores which I've used to derive at above conclusions

In [1]:

```

import pandas as pd
import numpy as np
from nltk import *
import matplotlib.pyplot as plt

from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, LabelEncoder
from sklearn.model_selection import train_test_split, cross_validate, learning_curve
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
from sklearn.naive_bayes import MultinomialNB, GaussianNB
from sklearn.metrics import precision_recall_fscore_support
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix

import seaborn as sns
sns.set(rc={'figure.figsize':(10,8)})
sns.set_style('white')
colors = sns.color_palette('pastel')

```

In [2]:

```

# Took reference from sklearn
# https://scikit-learn.org/stable/auto_examples/model_selection/plot_learning_curve.html

def plot_learning_curve(
    estimator,
    title,
    X,
    y,
    axes=None,
    ylim=None,
    cv=None,
    n_jobs=None,
    train_sizes=np.linspace(0.1, 1.0, 5),
):
    """
    Generate 3 plots: the test and training learning curve, the training
    samples vs fit times curve, the fit times vs score curve.

    Parameters
    -----
    estimator : estimator instance
        An estimator instance implementing `fit` and `predict` methods which
        will be cloned for each validation.

    title : str
        Title for the chart.

    X : array-like of shape (n_samples, n_features)
        Training vector, where ``n_samples`` is the number of samples and
        ``n_features`` is the number of features.
    """

```

```

y : array-like of shape (n_samples) or (n_samples, n_features)
    Target relative to ``X`` for classification or regression;
    None for unsupervised learning.

axes : array-like of shape (3,), default=None
    Axes to use for plotting the curves.

ylim : tuple of shape (2,), default=None
    Defines minimum and maximum y-values plotted, e.g. (ymin, ymax).

cv : int, cross-validation generator or an iterable, default=None
    Determines the cross-validation splitting strategy.
    Possible inputs for cv are:

    - None, to use the default 5-fold cross-validation,
    - integer, to specify the number of folds.
    - :term:`CV splitter`,
    - An iterable yielding (train, test) splits as arrays of indices.

    For integer/None inputs, if ``y`` is binary or multiclass,
    :class:`StratifiedKFold` used. If the estimator is not a classifier
    or if ``y`` is neither binary nor multiclass, :class:`KFold` is used.

    Refer :ref:`User Guide <cross_validation>` for the various
    cross-validators that can be used here.

n_jobs : int or None, default=None
    Number of jobs to run in parallel.
    ``None`` means 1 unless in a :obj:`joblib.parallel_backend` context.
    ``-1`` means using all processors. See :term:`Glossary <n_jobs>`
    for more details.

train_sizes : array-like of shape (n_ticks,)
    Relative or absolute numbers of training examples that will be used to
    generate the learning curve. If the ``dtype`` is float, it is regarded
    as a fraction of the maximum size of the training set (that is
    determined by the selected validation method), i.e. it has to be within
    (0, 1]. Otherwise it is interpreted as absolute sizes of the training
    sets. Note that for classification the number of samples usually have
    to be big enough to contain at least one sample from each class.
    (default: np.linspace(0.1, 1.0, 5))
"""
if axes is None:
    _, axes = plt.subplots(1, 3, figsize=(20, 5))

axes[0].set_title(title)
if ylim is not None:
    axes[0].set_ylim(*ylim)
axes[0].set_xlabel("Training examples")
axes[0].set_ylabel("Score")

train_sizes, train_scores, test_scores, fit_times, _ = learning_curve(
    estimator,

```

```

        X,
        Y,
        cv=cv,
        n_jobs=n_jobs,
        train_sizes=train_sizes,
        return_times=True,
    )
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    fit_times_mean = np.mean(fit_times, axis=1)
    fit_times_std = np.std(fit_times, axis=1)

    # Plot learning curve
    axes[0].grid()
    axes[0].fill_between(
        train_sizes,
        train_scores_mean - train_scores_std,
        train_scores_mean + train_scores_std,
        alpha=0.1,
        color="r",
    )
    axes[0].fill_between(
        train_sizes,
        test_scores_mean - test_scores_std,
        test_scores_mean + test_scores_std,
        alpha=0.1,
        color="g",
    )
    axes[0].plot(
        train_sizes, train_scores_mean, "o-", color="r", label="Training score"
    )
    axes[0].plot(
        train_sizes, test_scores_mean, "o-", color="g", label="Cross-validation score"
    )
    axes[0].legend(loc="best")

    # Plot n_samples vs fit_times
    axes[1].grid()
    axes[1].plot(train_sizes, fit_times_mean, "o-")
    axes[1].fill_between(
        train_sizes,
        fit_times_mean - fit_times_std,
        fit_times_mean + fit_times_std,
        alpha=0.1,
    )
    axes[1].set_xlabel("Training examples")
    axes[1].set_ylabel("fit_times")
    axes[1].set_title("Scalability of the model")

    # Plot fit_time vs score
    fit_time_argsort = fit_times_mean.argsort()

```



```

fit_time_sorted = fit_times_mean[fit_time_argsort]
test_scores_mean_sorted = test_scores_mean[fit_time_argsort]
test_scores_std_sorted = test_scores_std[fit_time_argsort]
axes[2].grid()
axes[2].plot(fit_time_sorted, test_scores_mean_sorted, "o-")
axes[2].fill_between(
    fit_time_sorted,
    test_scores_mean_sorted - test_scores_std_sorted,
    test_scores_mean_sorted + test_scores_std_sorted,
    alpha=0.1,
)
axes[2].set_xlabel("fit_times")
axes[2].set_ylabel("Score")
axes[2].set_title("Performance of the model")

return plt

```

```

In [3]: data = pd.read_csv('data.csv')
        data.head(3)

```

```

Out[3]:

```

	Description	Belts	Personal Injury	Property Damage	Commercial License	Commercial Vehicle	State	VehicleType	Year
0	'DISPLAYING EXPIRED REGISTRATION PLATE ISSUED ...	No	No	No	No	No	NC	'02 - Automobile'	201
1	'DRIVER FAIL TO STOP AT RED TRAFFIC SIGNAL BEF...	No	No	No	No	No	MD	'02 - Automobile'	201
2	'DRIVING UNDER THE INFLUENCE OF ALCOHOL PER SE'	No	No	No	No	No	MD	'02 - Automobile'	200

```

In [4]: data.info()

```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 70340 entries, 0 to 70339
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Description                          70340 non-null  object
1   Belts                               70340 non-null  object
2   Personal Injury                     70340 non-null  object
3   Property Damage                    70340 non-null  object
4   Commercial License                 70340 non-null  object
5   Commercial Vehicle                 70340 non-null  object
6   State                              70340 non-null  object
7   VehicleType                        70340 non-null  object
8   Year                                70340 non-null  object
9   Make                                70340 non-null  object
10  Model                               70340 non-null  object
11  Color                               70340 non-null  object
12  Contributed To Accident            70340 non-null  object
13  Driver Race                        70340 non-null  object
14  Gender                             70340 non-null  object
15  Driver City                        70340 non-null  object
16  Drive State                        70340 non-null  object
17  Violation Type                     70340 non-null  object
dtypes: object(18)
memory usage: 9.7+ MB
```

Feature Engineering

```
In [5]: desc_gp = data.groupby('Violation Type').agg({'Description':lambda x:' '.join
desc_gp.head()
```

```
Out[5]:
```

	Description
Violation Type	
Citation	'DISPLAYING EXPIRED REGISTRATION PLATE ISSUED ...
SERO	'STOP LIGHTS' Headlights 'STOP LIGHTS' 'Wheels...
Warning	'FAILURE OF MV OPER TO PRESENT EVIDENCE OF REQ...

```
In [6]: def get_counts(n, text):
text = text.replace('\n', ' ')
tokens = word_tokenize(text)
text = Text(tokens)
counts = {}

for i in range(1, n+1):
    counts[i] = FreqDist(ngrams(text, i))

return counts
```

```
In [7]: cit_c = get_counts(4, desc_gp.loc['Citation'].values[0])
        sero_c = get_counts(4, desc_gp.loc['SERO'].values[0])
        war_c = get_counts(4, desc_gp.loc['Warning'].values[0])
```

```
In [8]: cit_c[2].most_common(5)
```

```
Out[8]: [ (('ON', 'HIGHWAY'), 7897),
          (('FAILURE', 'TO'), 6171),
          (('VEHICLE', 'ON'), 5997),
          (('MOTOR', 'VEHICLE'), 5454),
          (('MPH', 'IN'), 4654)]
```

```
In [9]: cit_c[4].most_common(10)
```

```
Out[9]: [ (('MPH', 'IN', 'A', 'POSTED'), 4644),
          (('MOTOR', 'VEHICLE', 'ON', 'HIGHWAY'), 3935),
          (('PERSON', 'DRIVING', 'MOTOR', 'VEHICLE'), 3007),
          (('DRIVING', 'MOTOR', 'VEHICLE', 'ON'), 2880),
          (('VEHICLE', 'ON', 'HIGHWAY', 'WITHOUT'), 2096),
          (('PUBLIC', 'USE', 'PROPERTY', 'ON'), 1994),
          (('VEHICLE', 'ON', 'HIGHWAY', 'OR'), 1975),
          (('ON', 'HIGHWAY', 'OR', 'PUBLIC'), 1975),
          (('HIGHWAY', 'OR', 'PUBLIC', 'USE'), 1975),
          (('OR', 'PUBLIC', 'USE', 'PROPERTY'), 1975)]
```

```
In [10]: sero_c[2].most_common(5)
```

```
Out[10]: [ (('STOP', 'LIGHTS'), 991),
           (('Stop', 'Lights'), 406),
           (('LIGHTS', 'STOP'), 340),
           (('WINDOW', 'TINT'), 314),
           (('TAG', 'LIGHTS'), 206)]
```

```
In [11]: sero_c[4].most_common(5)
```

```
Out[11]: [ (('STOP', 'LIGHTS', 'STOP', 'LIGHTS'), 277),
           (('STOP', 'LIGHTS', 'Stop', 'Lights'), 129),
           (('Stop', 'Lights', 'STOP', 'LIGHTS'), 110),
           (('LIGHTS', 'STOP', 'LIGHTS', 'STOP'), 98),
           (('WINDOW', 'TINT', 'STOP', 'LIGHTS'), 92)]
```

```
In [12]: war_c[2].most_common(5)
```

```
Out[12]: [ (('FAILURE', 'TO'), 12498),
           (('DRIVER', 'FAILURE'), 6194),
           (('TO', 'DISPLAY'), 4713),
           (('SPEED', 'LIMIT'), 4415),
           (('LIMIT', 'OF'), 4306)]
```

```
In [13]: war_c[4].most_common(10)
```

```
Out[13]: [ (('EXCEEDING', 'THE', 'POSTED', 'SPEED'), 4294),
  (('THE', 'POSTED', 'SPEED', 'LIMIT'), 4294),
  (('POSTED', 'SPEED', 'LIMIT', 'OF'), 4294),
  (('DRIVER', 'FAILURE', 'TO', 'OBEY'), 4191),
  (('FAILURE', 'TO', 'OBEY', 'PROPERLY'), 3999),
  (('TO', 'OBEY', 'PROPERLY', 'PLACED'), 3999),
  (('OBEY', 'PROPERLY', 'PLACED', 'TRAFFIC'), 3999),
  (('PROPERLY', 'PLACED', 'TRAFFIC', 'CONTROL'), 3999),
  (('PLACED', 'TRAFFIC', 'CONTROL', 'DEVICE'), 3999),
  (('TRAFFIC', 'CONTROL', 'DEVICE', 'INSTRUCTIONS'), 3999)]
```

Using text classification on decription column to predict class and use that as a feature

```
In [14]: mp = {
    'Citation': 0,
    'SERO': 1,
    'Warning': 2,
}

rev_mp = {
    0: 'Citation',
    1: 'SERO',
    2: 'Warning',
}

X_train, X_test, y_train, y_test = train_test_split(data['Description'], data
count_vect = CountVectorizer(max_features=1000, ngram_range=(1, 3), analyzer
tfidf_transformer = TfidfTransformer()
X_train_counts = count_vect.fit_transform(X_train)
X_train_tfidf = tfidf_transformer.fit_transform(X_train_counts)
clf = MultinomialNB().fit(X_train_tfidf, y_train)


data['Description Pred'] = data['Description'].apply(lambda x: mp[clf.predict
```

```
In [15]: X_test_counts = count_vect.fit_transform(X_test)
X_test_tfidf = tfidf_transformer.fit_transform(X_test_counts)
clf.score(X_test_counts, y_test)
```

```
Out[15]: 0.7102644299118567
```

Exploratory Data Analysis


```
In [16]: plt.figure(figsize=(4, 5))
_ = data['Violation Type'].value_counts().plot.bar()
```



```
In [17]: data['Driver City'].value_counts()
```

```
Out[17]: 'SILVER SPRING'      17691
          GAITHERSBURG      7224
          GERMANTOWN        5838
          ROCKVILLE        5804
          WASHINGTON        2133
          ...
          SPRING             1
          'W HYATTSVILLE'  1
          'SN LUIS OBISP'    1
          GRUNDY             1
          MEQUON             1
          Name: Driver City, Length: 1890, dtype: int64
```

```
In [18]: _ = data['Year'].value_counts().head(40).plot.bar()
```



```
In [19]: data['Model'].value_counts()
```

```
Out[19]: 4S          7780
TK          4604
ACCORD      2696
CIVIC       2392
CAMRY       2375
...
'GRAN MARQ'      1
'MILLENNIA 4S'   1
'SLK 230'        1
LS450           1
CONVERT         1
Name: Model, Length: 3828, dtype: int64
```

Data Preprocessing

In [20]:

```
df = data.copy()
ordinal_enc = OrdinalEncoder()

cols = ['Description', 'Belts', 'Personal Injury', 'Property Damage',
        'Commercial License', 'Commercial Vehicle', 'State', 'VehicleType',
        'Year', 'Make', 'Model', 'Color', 'Contributed To Accident',
        'Driver Race', 'Gender', 'Driver City', 'Drive State', 'Violation Type']

df[cols] = ordinal_enc.fit_transform(data[cols])

df.head()
```

Out[20]:

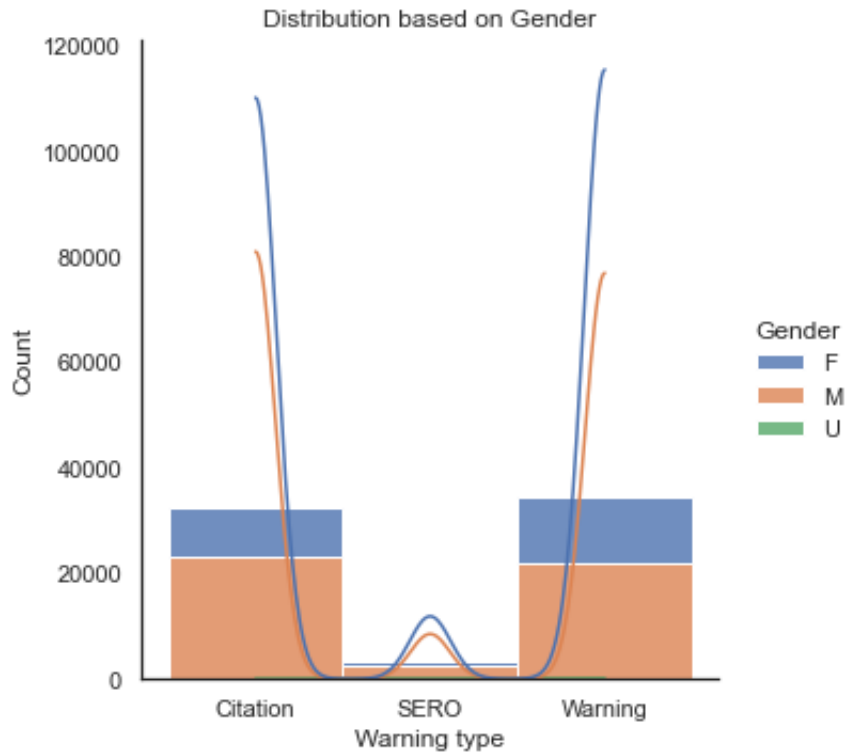
	Description	Belts	Personal Injury	Property Damage	Commercial License	Commercial Vehicle	State	VehicleType	Year
0	90.0	0.0	0.0	0.0	0.0	0.0	29.0	1.0	64.0
1	150.0	0.0	0.0	0.0	0.0	0.0	22.0	1.0	66.0
2	373.0	0.0	0.0	0.0	0.0	0.0	22.0	1.0	51.0
3	1799.0	0.0	0.0	0.0	0.0	0.0	22.0	1.0	63.0
4	90.0	0.0	0.0	0.0	1.0	0.0	22.0	1.0	61.0

Visualizations


In [21]:

```
_ = sns.displot(
    x = 'Violation Type',
    hue = 'Gender',
    kde = True,
    data = data,
    multiple = 'stack',
    alpha = 0.8,
)
plt.title('Distribution based on Gender')
plt.xlabel('Warning type')
plt.ylabel('Count')
```


Out[21]: Text(-7.27311840277776, 0.5, 'Count')




```
In [22]: plt.figure(figsize=(15, 5))
_ = sns.countplot(
    x = 'Drive State',
    hue = 'Violation Type',
    data = data,
    alpha = 0.8,
)
_ = plt.xticks(rotation=45)
_ = plt.legend(loc=1)
```




```
In [23]: plt.figure(figsize=(15, 5))
_ = sns.countplot(
    x = 'Color',
    data = data,
    alpha = 0.8,
)
_ = plt.xticks(rotation=45)
```

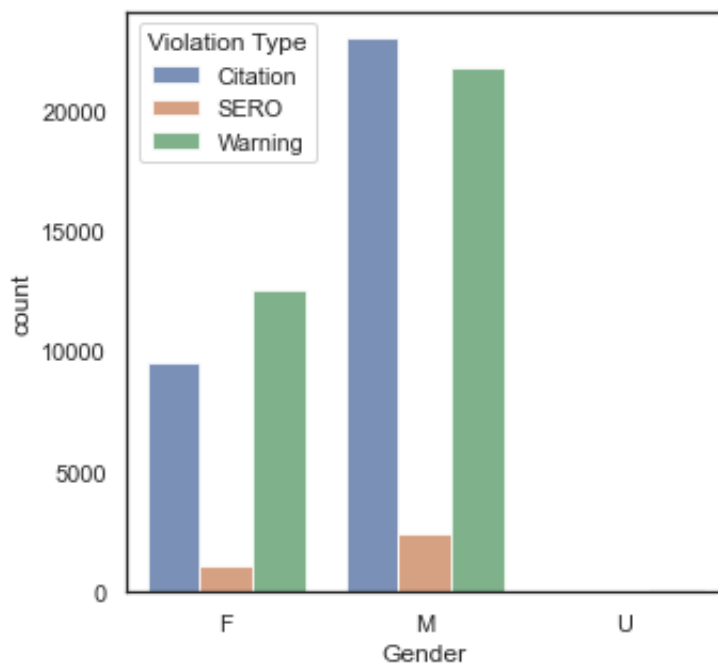


```
In [24]: plt.figure(figsize=(15, 5))
_ = sns.countplot(
    hue = 'Color',
    x = 'Violation Type',
    data = data,
    alpha = 0.8,
)
_ = plt.legend(loc=1)
```




Females have got lesser Citations than Males

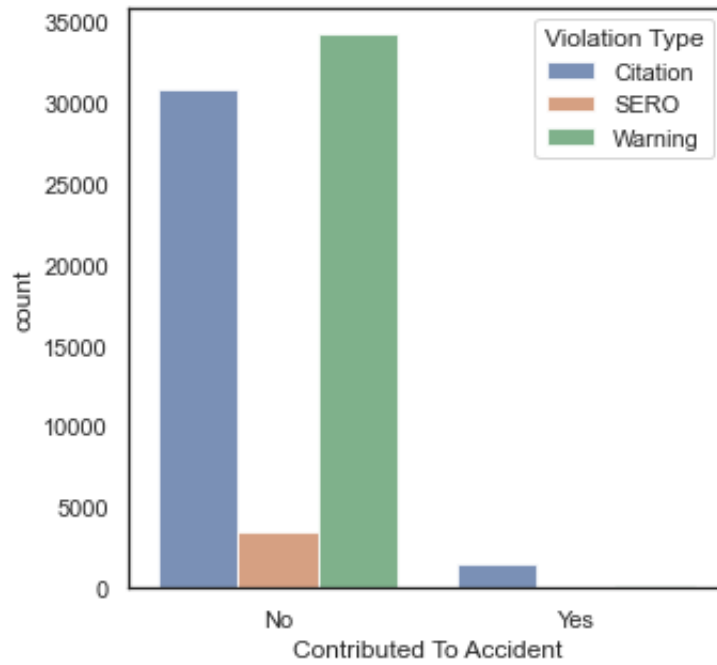
```
In [35]: plt.figure(figsize=(5, 5))
_ = sns.countplot(
    x = 'Gender',
    hue = 'Violation Type',
    data = data,
    alpha = 0.8,
)
```



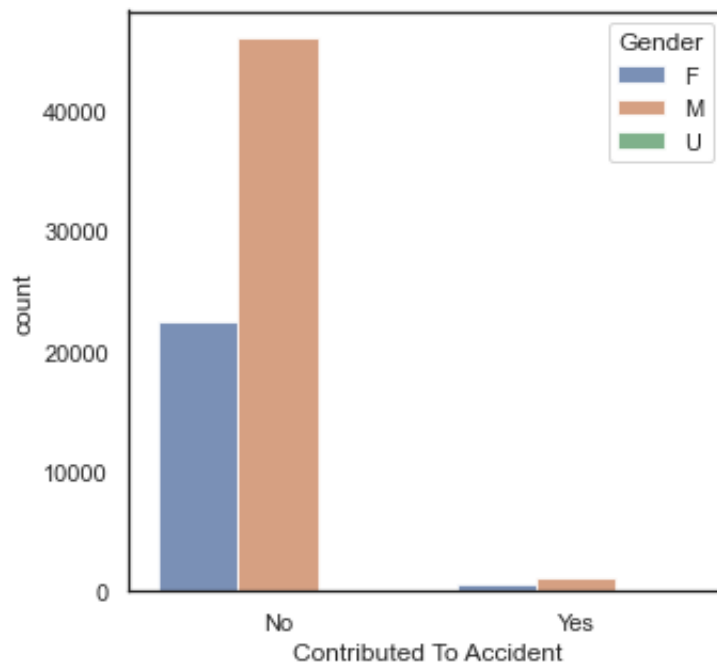
```
In [26]: plt.figure(figsize=(5, 5))
_ = sns.countplot(
    x = 'Driver Race',
    hue = 'Violation Type',
    data = data,
    alpha = 0.8,
)
_ = plt.xticks(rotation=45)
```



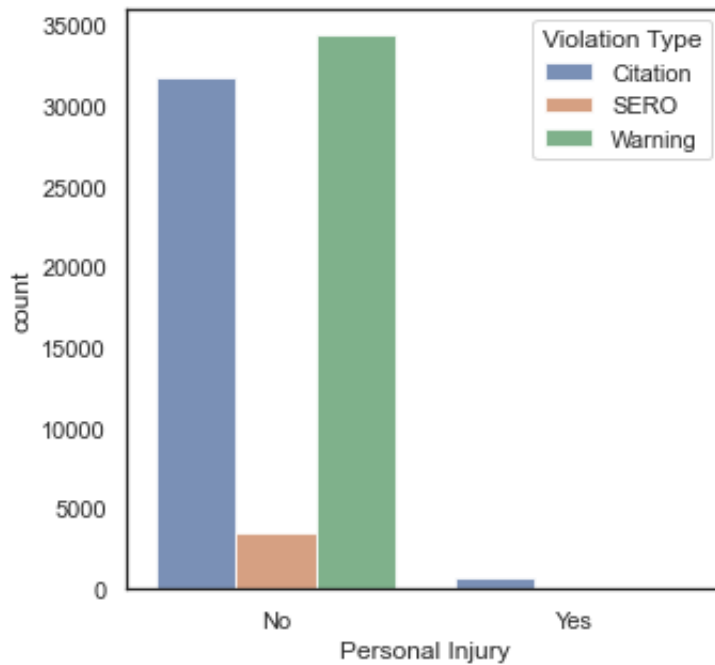
```
In [39]: plt.figure(figsize=(5, 5))
_ = sns.countplot(
    x = 'Contributed To Accident',
    hue = 'Violation Type',
    data = data,
    alpha = 0.8,
)
```



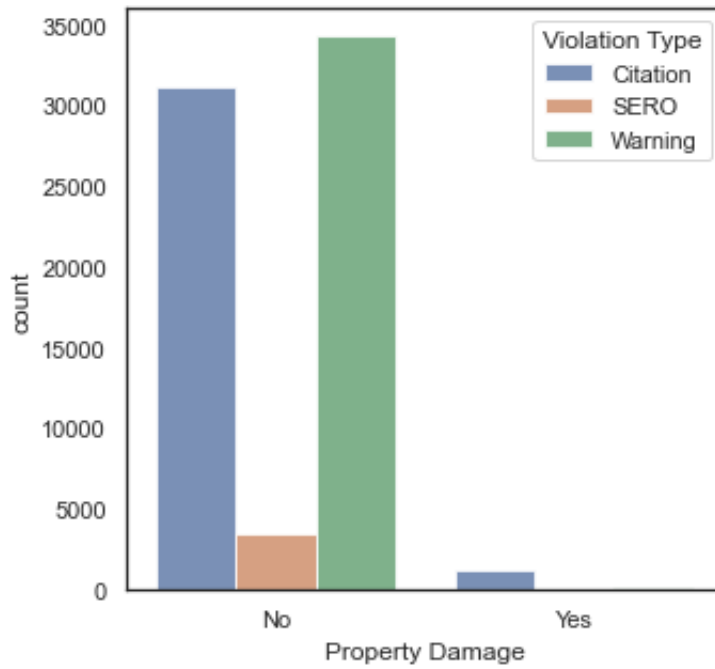
```
In [38]: plt.figure(figsize=(5, 5))
_ = sns.countplot(
    x = 'Contributed To Accident',
    hue = 'Gender',
    data = data,
    alpha = 0.8,
)
```



```
In [41]: plt.figure(figsize=(5, 5))
_ = sns.countplot(
    x = 'Personal Injury',
    hue = 'Violation Type',
    data = data,
    alpha = 0.8,
)
```




```
In [42]: plt.figure(figsize=(5, 5))
_ = sns.countplot(
    x = 'Property Damage',
    hue = 'Violation Type',
    data = data,
    alpha = 0.8,
)
```




```
In [27]: _ = plt.figure(figsize=(5, 5))
_ = sns.displot(
    x = 'Violation Type',
    hue = 'Contributed To Accident',
    kde = True,
    data = data,
    multiple = 'stack',
    alpha = 0.8,
)
_ = plt.title('Distribution based on Gender')
_ = plt.xlabel('Warning type')
_ = plt.ylabel('Count')
```

<Figure size 360x360 with 0 Axes>




In [28]:


```
= sns.scatterplot(  
    x = 'Drive State',  
    y = 'Gender',  
    hue = 'Violation Type',  
    data = df,  
)
```




```
In [29]: plt.figure(figsize=(5, 5))
_ = sns.barplot(
    x = 'Gender',
    y = 'Violation Type',
    hue = 'Contributed To Accident',
    data = df,
)
```

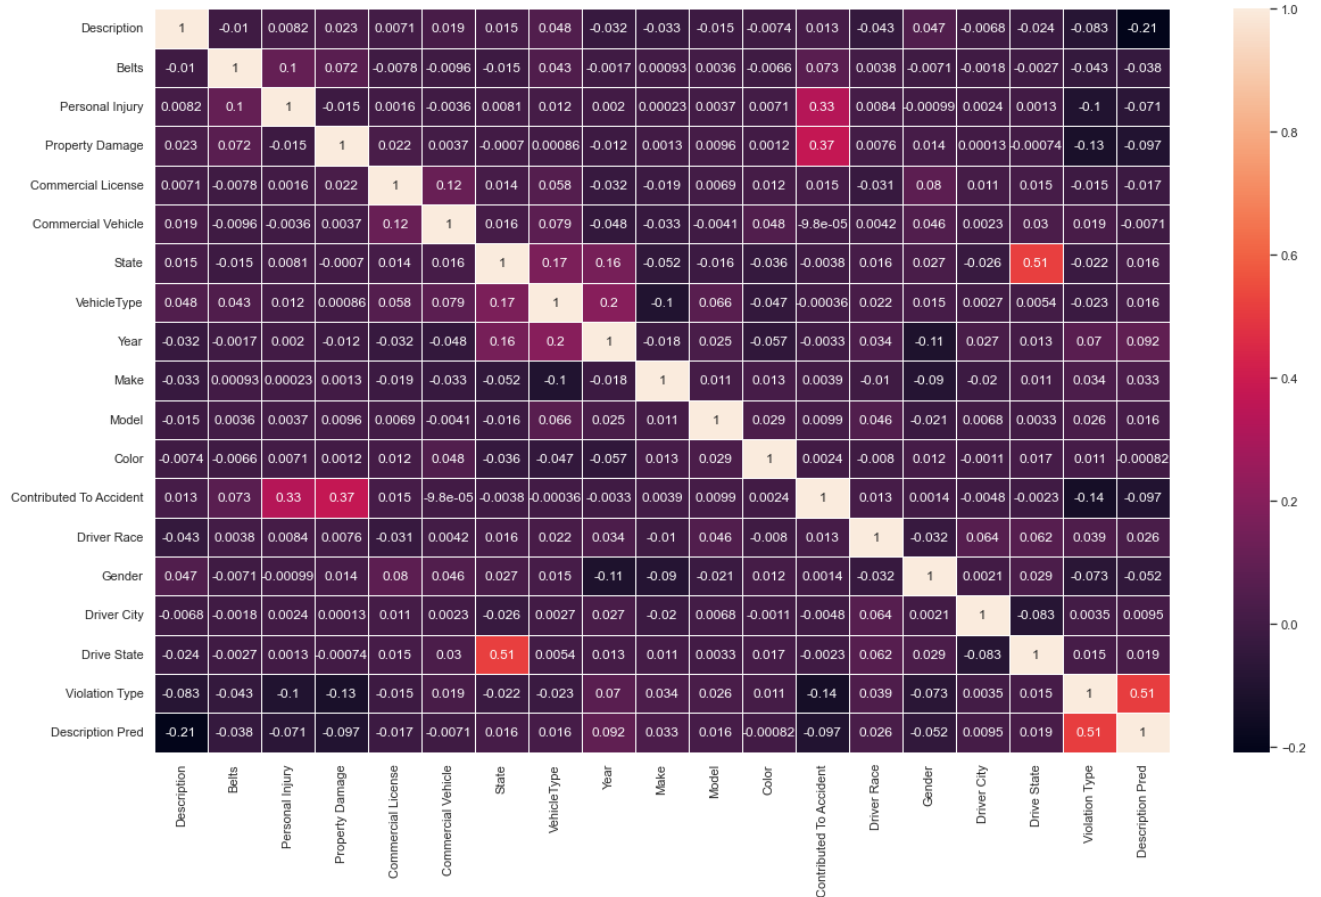
```
In [30]: _ = sns.scatterplot(  
    y = 'State',  
    x = 'Driver Race',  
    hue = 'Violation Type',  
    data = df,  
)
```



```
In [31]: plt.figure(figsize=(5, 5))
_ = sns.barplot(
    hue = 'Violation Type',
    y = 'Personal Injury',
    x = 'Gender',
    data = df,
)
```



```
In [43]: plt.figure(figsize = (20, 12))
corr = df.corr()
sns.heatmap(corr, annot = True, linewidths = 1)
plt.show()
```



```
In [123... df.corr()['Violation Type'].sort_values()
```

```
Out[123... Contributed To Accident    -0.139049
Property Damage              -0.125448
Personal Injury              -0.101189
Description                  -0.082798
Gender                       -0.072796
Belts                        -0.042997
VehicleType                  -0.023455
State                        -0.021977
Commercial License           -0.015411
Driver City                  0.003515
Color                        0.010616
Drive State                  0.014618
Commercial Vehicle           0.018542
Model                        0.025527
Make                         0.033530
Driver Race                  0.039167
Year                         0.070488
Description Pred              0.513349
Violation Type                1.000000
Name: Violation Type, dtype: float64
```

Modelling

In [206...

```

target_col = 'Violation Type'
target = df[target_col]
select_cols = [
    'Contributed To Accident',
    'Description',
    'Belts',
    'Property Damage',
    'Personal Injury',
    'Description Pred',
]
X_train, x_test, y_train, y_test = train_test_split(df[select_cols], target,

```

Random Forest Classifier

In [207...

```

forest = RandomForestClassifier(n_jobs=-1)
forest.fit(X_train, y_train)
forest_y = forest.predict(x_test)

test_score = forest.score(x_test, y_test)
print(f"Accuracy of the RandomForestClassifier: {test_score:.4f}")

```

Accuracy of the RandomForestClassifier: 0.7912

In [112...

```

cv_results = cross_validate(forest, df[select_cols], target)
scores = cv_results["test_score"]
print(f"Accuracy score via cross-validation:\n"
      f"{scores.mean():.3f} +/- {scores.std():.3f}")

```

Accuracy score via cross-validation:
0.790 +/- 0.003

In [113...

```
print(classification_report(y_test, forest_y, target_names=['Citation', 'SERO

```

	precision	recall	f1-score	support
Citation	0.86	0.65	0.74	6456
SERO	0.99	1.00	0.99	694
Warning	0.74	0.90	0.81	6918
accuracy			0.79	14068
macro avg	0.86	0.85	0.85	14068
weighted avg	0.81	0.79	0.79	14068

In [212...

```

res = []
for c in [0, 1, 2]:
    prec, recall, _, _ = precision_recall_fscore_support(
        np.array(y_test) == c,
        np.array(forest_y) == c,
        pos_label=True,
        average=None,
    )
    res.append([rev_mp[c], recall[0], recall[1]])

pd.DataFrame(res, columns = ['class', 'sensitivity', 'specificity'])

```

Out[212...


	class	sensitivity	specificity
0	Citation	0.905281	0.657993
1	SERO	0.999477	0.988489
2	Warning	0.690952	0.895619

In [40]:

```

_ = plot_learning_curve(forest, 'Random Forest Model', X_train, y_train)

```



KNN

In [208...

```

knn = KNeighborsClassifier(100)
knn.fit(X_train, y_train)
knn_y = knn.predict(x_test)

test_score = knn.score(x_test, y_test)
print(f"Accuracy of the KNN: {test_score:.2f}")

```

Accuracy of the KNN: 0.78

```
In [209... cv_results = cross_validate(knn, df[select_cols], target)
scores = cv_results["test_score"]
print(f"Accuracy score via cross-validation:\n"
      f"{scores.mean():.3f} +/- {scores.std():.3f}")
```

Accuracy score via cross-validation:
0.641 +/- 0.003

```
In [210... print(classification_report(y_test, knn_y, target_names=['Citation', 'SERO',
```


	precision	recall	f1-score	support
Citation	0.85	0.63	0.73	6456
SERO	0.92	0.94	0.93	695
Warning	0.72	0.89	0.80	6917
accuracy			0.78	14068
macro avg	0.83	0.82	0.82	14068
weighted avg	0.79	0.78	0.77	14068

```
In [211... res = []
for c in [0, 1, 2]:
    prec, recall, _, _ = precision_recall_fscore_support(
        np.array(y_test) == c,
        np.array(knn_y) == c,
        pos_label=True,
        average=None,
    )
    res.append([rev_mp[c], recall[0], recall[1]])

pd.DataFrame(res, columns = ['class', 'sensitivity', 'specificity'])
```

```
Out[211... class sensitivity specificity
0 Citation 0.905413 0.632900
1 SERO 0.996037 0.939568
2 Warning 0.667319 0.893017
```

```
In [45]: _ = plot_learning_curve(knn, 'KNN Model', X_train, y_train)
```



SVM

In [46]:

```
svm = SVC()
svm.fit(X_train, y_train)
svm_y = svm.predict(x_test)

test_score = svm.score(x_test, y_test)
print(f"Accuracy of the SVM: {test_score:.2f}")
```

Accuracy of the SVM: 0.63

In [47]:

```
cv_results = cross_validate(svm, df[select_cols], target)
scores = cv_results["test_score"]
print(f"Accuracy score via cross-validation:\n"
      f"{scores.mean():.3f} +/- {scores.std():.3f}")
```

Accuracy score via cross-validation:
0.634 +/- 0.002

In [48]:

```
print(classification_report(y_test, svm_y, target_names=['Citation', 'SERO',
```

	precision	recall	f1-score	support
Citation	0.67	0.47	0.55	6495
SERO	0.68	0.94	0.79	709
Warning	0.61	0.76	0.67	6864
accuracy			0.63	14068
macro avg	0.66	0.72	0.67	14068
weighted avg	0.64	0.63	0.63	14068


```

In [49]:
res = []
for c in [0, 1, 2]:
    prec, recall, _, _ = precision_recall_fscore_support(
        np.array(y_test) == c,
        np.array(svm_y) == c,
        pos_label=True,
        average=None,
    )
    res.append([rev_mp[c], recall[0], recall[1]])

pd.DataFrame(res, columns = ['class', 'sensitivity', 'specificity'])

```

```

Out[49]:


```

	class	sensitivity	specificity
0	Citation	0.804833	0.471286
1	SERO	0.977094	0.935120
2	Warning	0.534148	0.758159

```

In [50]:
_ = plot_learning_curve(svm, 'SVM Model', X_train, y_train)

```



Part - C

"Exascale Computing and Big Data" - by Daniel A. Reed and Jack Dongarra

Motivation

The paper begins by discussing some of the major challenges that are currently being solved by computational models, such as climate modeling and predicting black hole behavior, and then goes on to discuss how solving these problems require some investment in advanced computing research, design, and global collaboration.

Scientific and engineering opportunities with the rise of big data analytics

The authors begin by discussing how big data analytical systems were critical in the discovery of the Higgs boson particle, and they go on to discuss how various fields other than physical sciences are leveraging the potential of big data. It has been reported that some astronomers are now frequently querying existing datasets to discover previously unknown patterns and trends, which is very interesting to know. Big data is also playing an important role in the healthcare domain, with many of these techniques being increasingly used to conduct comparative and longitudinal analyses of cancer treatment regimes. The author also discussed some lesser-known fields where big data analytics is being used; currently, semantic graph visualization tools and recommender systems are being used to identify relevant topics and suggest relevant papers for study.

Technical challenges towards high-end computing

The study explored a wide range of technological obstacles, including software, hardware, and scaling challenges. The author references a survey undertaken by the US Department of Energy, which claims that new designs and technologies are required to reduce the energy requirement to a more manageable and economically reasonable level in order to satisfy the present compute power demand.

Outlining some of the strategies that can address these issues

- Locality-aware algorithms (MapReduce toolkits) and software will be needed to maximize computation performance and reduce energy needs
- New memory technologies, including processor in-memory, stacked memory, and nonvolatile memory approaches are needed to minimize data movement and energy use
- More expressive programming methods are required to deal with the parallelism and frequent errors of exascale computing systems.
- Given the need for systemic resilience in the face of component failures, systems should incorporate a wide range of resilience mechanisms, such as geo-distribution, automated restart and failover, failure injection, and introspective monitoring.